

Week Plan 2026-07-06

Week Plan - Mon 6 to Fri 10 July 2026

Five working days (UAE Mon-Fri).

Roughly 3.5 days of estimated build work plus admin, so the week is full but achievable if the order holds. The one hard dependency: the humaniser fix comes before any new SEO content ships, because publishing through a failing process just adds pages you will have to redo.

The five workstreams

Workst Est. Gates

ream effort /

depend
s on

1 Human ~0.5 Blocks

ser day SEO

solutio researc content

n (AI- h +

detecti setup publishi

on fix) ng

2 SEO ~0.5 Needs

tidy-up day humani

finish ser

(compa done

re + before

lead- content

gen)

resume
s

3 Routine ~0.5 None

day

prompt admin

+

CLAUD

E.md

rewrite

(learn-

from-

builds)

4 Atlas ~1.5 None,

build days but

finish + heavies

test, t
integrat
e
Moonie'
s plan

5 Orwell ~1.5 Renam

Tours days e
build decisio
(renam n
e first) needed

up front

That totals ~4.5 days of work against 5 available. Atlas and Tours cannot both fully complete in one week alongside the rest. The plan front-loads the blockers and the quick wins, protects two clear days for Atlas, and starts Tours, with an honest note that Tours likely spills into the following week.

**Monday 6 July - Unblock
and clear the small stuff**

The theme today is removing blockers so the rest of the week runs clean. Humaniser first because it gates content, then the two half-day admin jobs.

Morning: Humaniser solution (workstream 1)

- ☐ Define the actual failure. Which detector is the content failing (Originality.ai, GPTZero, Copyleaks, Turnitin)? Get one real failing sample and its score before researching fixes.
- ☐ Research the current state of humanising. Note: your earlier position was to reframe around genuine voice quality rather than gaming detectors. That was right in principle but it is now failing in practice, so this pass needs a working process, not a philosophy. Both can be true: build real voice AND pass detection.

- ☐ Decide the approach: prompt-level changes to the writing routine (burstiness, sentence-length variance, first-person specifics, real data points), a post-generation humanising pass, or a third-party API in the pipeline. Weigh cost and reliability.
- ☐ Produce a testable spec: take one failing article, apply the new process, re-run the same detector, record the before/after score. No spec ships without a passing re-test.

Afternoon: Routine prompt + CLAUDE.md rewrite (workstream 3)

- ☐ Open Claude Code in each relevant repo (the daily-brief Brain repo, and the content-build repos for the sites).

- ☐ Paste the current routine prompts in and ask Claude Code to rewrite both the prompts and the CLAUDE.md files so the system learns from past builds and produces clean builds going forward. Frame it explicitly: "read the last N builds, find what went wrong or was redone, encode the fix as a standing rule."
- ☐ The architect-loop research just filed in the vault ([Architect Loop](#) if built, or the research note) is directly relevant here: the frozen-gates and repo-as-memory discipline is exactly what "learn from builds and build clean" means. Pull the useful rules from it into the CLAUDE.md rewrites.
- ☐ Commit the rewritten prompts and CLAUDE.md files. Note in each what changed and why.

End-of-day check: humaniser has a tested spec, routine prompts and CLAUDE.md files are rewritten and committed. Both blockers cleared.

Tuesday 7 July - Finish the SEO tidy-up

With the humaniser fixed, the content side of the tidy-up can resume. The technical fixes from the June indexing programme are mostly done; this is the finishing pass.

SEO tidy-up finish (workstream 2)

Grounding from the June indexing work, current known state:

- ☐ **RepatriationFuneral** was the worst site (~2,193 not-indexed, dominated by canonical duplication). Confirm the canonical and duplication repair landed and the not-indexed pile is shrinking in GSC. This was active-repair-first priority.

- ❑ **PetTransportGlobal** had 114 hard 404s and ~1,069 not-indexed despite being highest-traffic. Confirm the 404s are cleared and redirects are in place.
- ❑ **MortgageCompare** calculator pages: the five bank calculator pages were flagged as content-similarity duplicates (Google deduplicating near-identical pages). This needs bank-specific written content on each. Apply the now-fixed humaniser process here.
- ❑ **PestRemovalUSA**: sitemap-coverage and internal-linking fix (was ruled out as a build failure, 74 discovered-not-indexed). Confirm sitemap covers all live pages.
- ❑ **CloseProtectionHire**: confirm the robots.txt fix deployed (it had a bare User-agent line, no sitemap reference).

- ☐ Run the Monday cadence loop from the action plan as the closing check: connector health (GA4/GSC/Semrush green in Windsor), indexed count vs last week per site, sitemap count vs live pages.
- ☐ Mark the tidy-up formally done in the vault, or list exactly what remains and park it explicitly.

End-of-day check: portfolio is in a tidy, measurable state or the remainder is explicitly parked with reasons. Content process is now safe to run.

Wednesday 8 + Thursday 9 July - Atlas

Two protected days for the heaviest workstream. Atlas is the centrepiece agentic system (auto-research and auto-generate programmatic SEO sites across ~580,000 niche-geography cells).

Atlas build finish + test (workstream 4)

- ☐ Ground first: read the current Atlas repo state and the Global Lead Engine framework v1.0. Reconcile what the build docs claim against what actually exists in the repo (the architect-loop "trust git over the diary" rule applies).
- ☐ Integrate Moonie's plan. Feed Moonie's plan into Claude Code and have it integrate with the existing Atlas workflow. Before integrating: get Claude Code to read both the existing workflow and Moonie's plan, list the conflicts and overlaps, and propose the merge, rather than blind-merging.
- ☐ Finish the outstanding build items on the seed-agent architecture (seed to orchestrator to spawned workers, self-learning loop, memory).

- ☐ Test. This is the step most likely to get skipped under time pressure. Define what "Atlas works" means as a concrete test (e.g. seed produces one cell end-to-end, or the orchestrator spawns and a worker completes one task) and run it.
- ☐ Commit and write the Atlas state to its handoff doc so the next session resumes cleanly.

Note: if the humaniser or SEO work overran Monday/Tuesday, Atlas is where the pressure lands. Protect these two days. If something has to give, it is the start of Tours (Friday), not Atlas testing.

End-of-day Thursday check: Atlas build finished, Moonie's plan integrated, one concrete end-to-end test passed.

**Friday 10 July - Orwell
Tours: rename and begin**

Tours is a ~1.5 day job, so Friday starts it and it spills into next week. That is expected, not a failure. Today: settle the name and get the build moving.

Step 1: Rename (do this first, it touches everything)

The project is currently "Orwell Tours" and you want it renamed to something tech, real estate, or both. It is a white-label 3D property tour platform (phone walkarounds to Gaussian splat scenes to WhatsApp tour links) for Dubai agencies.

☐ Pick the name. Candidates spanning the three angles:

- Real-estate led: **Orwell Spaces, Orwell Estates, TourNest, Splat Estates**
- Tech led: **Orwell Render, SplatSpace, Orwell 3D, Volumetric**
- Both: **Orwell Spaces** or **Estate Render** (real estate + the tech verb)
- My pick if you want one: **Orwell Spaces**. It keeps the Orwell brand line consistent with the rest of the portfolio, reads as real estate, and "spaces" carries the 3D/volumetric sense without being gimmicky. Decide, do not agonise; the name matters less than shipping.

- Once chosen, rename before writing more code: the Neon project (currently orwell-tours), the R2 bucket, the Render services, the repo, and the app identifiers. Renaming later is far more painful.

Step 2: Begin the build

- Ground in the existing Orwell Tours blueprint (the mobile-app model: Expo/React Native agent app, Render + Neon + Cloudflare R2 + KIRI Engine for reconstruction Phase 0/1, RunPod for Phase 2).
- Work through the browser-prompt setup cards for the NOW-tier accounts (Neon, R2, Render) using the new name.
- Start the Expo app scaffold and the guided-capture flow.

- Accept this runs into next week.
Write the Tours state to a handoff doc so Monday resumes cleanly.

End-of-week check: name settled and applied, Tours infrastructure accounts set up, app scaffold started.

Honest read on the week

- **The week is full.** 4.5 days of work in 5 days with admin overhead. It holds only if Monday's blockers clear on time. If the humaniser research runs long, it eats into everything downstream.
- **The real risk is Atlas testing getting skipped.** It is the most valuable workstream and the easiest to half-finish. The plan protects two full days for it. Guard them.
- **Tours will not finish this week.** That is built into the plan. Friday is a start, not a completion. Do not let the desire to finish Tours pull time from Atlas testing.

- **Dependency reminder:** humaniser before SEO content, always. And rename Tours before building it.
- **Not in this plan but on the horizon:** the TOLATA hearing is 28 July. Not this week, but it will need prep time in the next fortnight. Flagged so it does not ambush you.

Open loops carried in

- Humaniser: earlier reframe was voice-quality over detector-gaming; now needs a process that does both.
- Atlas: Moonie's plan not yet integrated; self-learning loop and memory outstanding.
- Tours: rename undecided; infra accounts not yet created under final name.
- SEO: RepatriationFuneral canonical repair and PetTransport 404 clearance need GSC confirmation.
- Routine prompts and CLAUDE.md rewrites: not yet done.